



Assignment: Distributed Banking System

by

Name	Matric No.	Work %
Timothy Chang	<i>U2220136J</i>	33.33%
O Jing	<i>U2223350G</i>	33.33%
Neo Zhi Xuan	<i>U2222293E</i>	33.33%

SC4051 Distributed Systems

NANYANG TECHNOLOGICAL UNIVERSITY

April 2026

Contents

1	Introduction & System Overview	1
2	Message Design & Wire Protocol	1
2.1	Packet Header Format	1
2.2	Type-Length-Value (TLV) Payload Encoding	2
2.3	Callback Push Packet	3
2.4	Cross-Language Marshalling Challenges	3
3	System Architecture & Design	4
3.1	Server Architecture (Go)	4
3.2	Client Architecture (C++)	4
3.3	Concurrency Model for Accounts	5
4	Banking Services & Custom Operations	6
4.1	Core Services (S1–S4): Open, Close, Deposit, Withdraw	6
4.2	Monitor Service (S5)	6
4.3	GetBalance (Idempotent Operation)	7
4.4	Transfer (Non-Idempotent Operation)	7
5	Fault Tolerance & Invocation Semantics	8
5.1	At-Least-Once Semantics	8
5.2	At-Most-Once Semantics	8
5.3	Loss Simulator Design	9
5.4	Experimental Validation	10
5.5	Results & Discussion	10
6	Testing Strategy	11
6.1	Unit Tests (C++ Client)	11
6.2	Integration Tests (End-to-End)	11
7	Assumptions and Limitations	12
7.1	System Assumptions	12
7.2	Known Limitations	12
8	Conclusion	12

1 Introduction & System Overview

The project implements a distributed banking application, using a custom client-server architecture that uses UDP-based inter-process communication. The system is cross-language interoperable and consists of a C++ client application and a Go-based server. These components communicate via a custom binary protocol, without Remote Procedure Call (RPC) frameworks.




Figure 1: High-level architecture of the distributed banking system.

As shown in Figure 1, the client and server communicate via UDP datagrams, on known IP addresses and ports. The client constructs request messages and waits for responses, and may retry depending on the protocol. The server listens for requests and processes each one, before returning a response.

The client is written in C++ due to the language's ability to handle low level byte control for packet construction, as well as portable socket handling across Windows and POSIX. Meanwhile, the choice of Go for the server provides an in-depth networking library and built-in synchronisation primitives for concurrency, to allow for simple implementation of UDP sockets and concurrent client handling.

The presence of cross-language components necessitates handling of various constraints. Firstly, the data boundary and marshalling conventions for all data types must be aligned on both components. Next, it is important to standardise encodings so that both components interpret numeric bytes identically. Network byte order (big-endian) is used for all integer fields, whereas IEEE754 double format (64-bit) is used for floating-point balances.

This report analyses the unreliability of UDP, and the trade-offs between using different semantic modes to handle it, via explicit retry and duplicate-detection mechanisms.

2 Message Design & Wire Protocol

2.1 Packet Header Format

Each message sent between client and server begins with an 18-byte header, followed by a payload of variable length. It encodes information about protocol control, through these fields:

Field	Size	Description
Type	1 byte	Indicates whether the datagram is a request (0x00), a normal reply (0x01), or a server-initiated callback (0x02).
Semantics Flag	1 byte	Encodes invocation semantics: 0x00 for at-least-once, 0x01 for at-most-once.
Request ID	4 bytes	An identifier chosen by the client, used for duplicate detection and request-response matching.
Client IPv4 Address	4 bytes	The sender's IPv4 address. Explicitly informs the server about the client's identity.
Client Port	2 bytes	The UDP port number of the sender.
Status Code	2 bytes	Indicates success or error conditions in replies; non-zero values represent errors.
Content Length	4 bytes	Length (in bytes) of the payload that follows.




Figure 2: Packet structure and field descriptions for the request header.

This header is inspired by traditional RPC or network protocol headers. A fixed header size allows the client and server to parse data with a consistent offset. It also allows for robust validation of datagrams that are truncated. Simultaneously, all multi-byte fields are in big-endian format, aligning the client and server's method of interpreting numeric values.

2.2 Type-Length-Value (TLV) Payload Encoding

After the header, each message carries a TLV-encoded payload of application-specific fields. Fields are optional, and only serialised when it is present. Each field has this format:




Figure 3: Packet structure and field descriptions for the TLV encoding format.

For example, the request "OpenAccount (Alice, pass1234, SGD, 1000.00)" attempts to open an account, and might include these TLV fields:

Table 1: TLV fields for an OpenAccount request example.

Field Name	Length	Description
Account Holder Name	N bytes	UTF-8 encoded string representing the user's name. The length varies depending on the number of characters.
Password	Up to 8 bytes	Password represented as raw bytes. The client enforces a maximum length of 8 bytes. Shorter passwords are stored at their actual length.
Currency Type	1 byte	Encoded as a predefined enumeration value representing supported currencies (e.g., SGD, USD).
Initial Balance	8 bytes	IEEE-754 double precision floating-point number encoded in big-endian format.

The client packs each numeric or string value without additional padding, and the length specifies the exact number of bytes. For example, if the user "Alice" (5 letters) chooses password "pass1234" (8 chars), currency SGD (id=1), and balance 1000.00, the TLV encoding for that request would look like:

1	10 00 00 00 01 01	[ServiceID field: ID=16, Len=1, Value=0x01 (OpenAccount service)]
2	11 00 00 00 01 01	[MsgType field: ID=17, Len=1, Value=0x01 (Request)]
3	01 00 00 00 05 41 6C 69 63 65	[Name field: ID=1, Len=5, Value="Alice"]
4	02 00 00 00 08 70 61 73 73 31 32 33 34	[Password field: ID=2, Len=8, Value="pass1234"]
5	03 00 00 00 01 01	[Currency field: ID=3, Len=1, Value=0x01 (SGD)]
6	04 00 00 00 08 40 8F 40 00 00 00 00 00	[Balance field: ID=4, Len=8, Value=1000.00]

Listing 1: TLV Encoding Example

The server, upon receiving this TLV payload, unpacks each field by reading the ID and length, then extracts the number of bytes as the value. It interprets the raw bytes based on the field's type - converting an 8-byte double to Go's float64 in big-endian, or treating password bytes as an ASCII string.

Meanwhile, certain fields are fixed-size. In the Go server, an 8-byte password field is enforced. This will be discussed in detail in Section 2.4.

Hence, this format handles data cleanly. The use of ID and length allows for robust parsing. While the encoding is handcrafted in code, it creates an effective platform-independent message format.

2.3 Callback Push Packet

Within the lab requirements, the server should "push" updates to all registered clients during a monitoring interval. When an account is opened, closed, deposited, or withdrawn, the server triggers a callback packet to the client. It uses a different packet format:

Field Name	Size	Description
MsgType	1 byte	Always set to 0x02 to differentiate callback packets from normal replies.
Service ID	1 byte	Indicates the operation to trigger (e.g., OpenAccount).
Account Number	4 bytes	Integer account number (big-endian).
Name Length	4 bytes	Length of the account holder's name in bytes.
Name	Variable	UTF-8 encoded string representing the account holder's name.
Currency Type	1 byte	The account's currency code as an enumeration.
Balance	8 bytes	IEEE-754 double precision floating-point value (big-endian) representing the updated balance.




Figure 4: Packet structure and field descriptions for the callback message.

For example, if account 12345 (“Bob”) with currency USD (0x00) now has balance 250.75, the server might send:

```

1  02 03          [MsgType=0x02, ServiceID=0x03 (Deposit triggered update)]
2  00 00 30 39    [Account#=12345 (big-endian)]
3  00 00 00 03 42 6F 62 [NameLen=3, Name="Bob"]
4  00            [Currency=0x00 (USD)]
5  40 6F 40 00 00 00 00 00 [Balance=250.75 (IEEE-754)]

```

Listing 2: Callback Packet Example

The client's UDP socket listens for these packets and, upon receipt, displays the updated account information to the user.

The benefits of having a fixed, known structure is that it avoids the overhead that TLV packets face with self-describing fields. TLV packets require 5 additional bytes for each field to store FieldID and Length, making it slower to encode, decode, and process.

2.4 Cross-Language Marshalling Challenges

Implementing the client and server in different languages always leads to challenges in data coordination, due to the differences in how each language defines its data types and mechanisms. The first challenge was designing the password field. The Go server defines passwords as a fixed 8-byte array, whereas the C++ client uses 'std::string'. To overcome this, the protocol's TLV password field was fixed to 8 bytes on the wire. The client enforces that the password string does not exceed 8 bytes, rejecting or truncating longer inputs.

The second issue is the different representation of floating-point values in memory across both languages, where C++ uses 'double' and Go uses 'float64'. While both languages use IEEE-754 64-bit binary representation for these numbers, they could differ in endianness, and lead to incorrect decoding of raw bytes. Hence, for all floating-point transmissions, values are converted into IEEE-754 format encoded in big-endian order before sending, and recipients would reverse the process. This ensures the same byte sequence is interpreted reliably and correctly.

3 System Architecture & Design

3.1 Server Architecture (Go)

The Go server is implemented with a layered design, illustrated in Figure 5.




Figure 5: Go server architecture.

The core components of the server architecture are divided as follows:

Main & Dispatcher:

The main function sets up a UDP socket and constructs the server components. The dispatcher then reads incoming UDP packets, unmarshals the payload and header, before routing each request to an appropriate service handler. Its behaviour depends on the chosen invocation semantics: for *at-least-once* mode, it forwards each request to the handler, while in *at-most-once* mode, it checks the request against a history table to eliminate duplicate requests.

Handler:

The handler receives a request from the dispatcher and chooses a service to trigger. It contains the application logic for services like opening, closing, depositing, and withdrawing from an account. It decodes the TLV fields from the payload and performs validation before invoking the appropriate service function and encoding a TLV response.

Banking Service & Memory Store:

The banking service and memory store layers handle the data operations. The service layer provides operations like `OpenAccount` that directly manipulate the storage. The memory store maintains a map of accounts and uses synchronisation primitives to ensure that concurrent access to data is safe and does not lead to race conditions.

Monitor Manager:

The monitoring feature is implemented within the banking service with a publish-subscribe mechanism. It stores a map of registered clients with their defined expiration times and handles the goroutines for dispatching account updates using callback packets.

By separating the logic into layers, we can isolate concerns. For example, semantics logic is handled in the dispatcher. This does not clutter banking service logic. Protocol encoding is handled only once in the dispatcher and handler. This allows dependency injection in tests and mocking of the UDP socket at the top level.

The architecture is wired by a `BuildHandler` function that exposes the banking service and monitor manager to all handlers. This ensures that the dispatcher can make calls to the handler without having to know the underlying banking or monitoring details.

3.2 Client Architecture (C++)

The C++ client is implemented with a three-layer architecture.




Figure 6: C++ client layered architecture.

When sending a request from the application layer, the client runs an interactive looping Command Line Interface (CLI) that takes a command from the user corresponding to banking operations. It communicates this request to the protocol layer. The protocol layer constructs a request message by populating the 18-byte header with the chosen invocation semantics flag and a fresh or zeroed request ID. It then converts numbers into big-endian format before encoding arguments into a TLV payload for cross-platform consistency. This layer also manages the decoding of replies and callbacks into structured data. Finally, the network layer abstracts UDP communication by providing send and receive operations over sockets.

When the client receives a response, the response takes the reverse path. It is received at the network layer. Then, the protocol layer validates headers and parses the payload, and thereafter the application layer presents the results or callback updates to the user. This design ensures that user logic, message formatting, and network communication are cleanly decoupled and easy to maintain.

At the same time, dependency injection was used to provide the client with access to various services. The client was able to access interfaces like `BaseSocket` and `BaseCommandEncoder` during its lifecycle. This allows the system to mock components for unit testing, and reduces coupling between the interface and networking/marshalling internals.

The dependency injection wiring at the client's entry point:

Listing 3: Constructor-based dependency injection in `main.cpp`

```
// Create concrete implementations
auto udpSocket = std::make_unique<NetworkUtils::UDPSocket>(serverIp, serverPort);
auto cmdEncoder = std::make_unique<Protocol::CommandEncoder>();
auto msgSerializer = std::make_unique<Protocol::MessageSerializer>();

// Inject via constructor
auto bankClient = std::make_unique<BankClient>(
    std::move(bankIO), std::move(udpSocket),
    std::move(cmdEncoder), std::move(msgSerializer), flag);
```

As `BankClient` accepts `BaseSocket` and `BaseCommandEncoder` interfaces, the test suite will substitute `MockSocket` and `MockEncoder` to verify the request assembly.

3.3 Concurrency Model for Accounts

While the report permits a single-threaded server, our system was designed to allow multiple clients to operate concurrently on the server, and requires various synchronisation strategies to align this.

The in-memory account map, where account numbers were mapped to account data, is protected by a global read-write mutex which synchronises account lifecycle operations like `CreateAccount` and `GetAccount`. While multiple read queries can be run under a shared read lock, write-based operations utilise an exclusive write lock.

Each account's state is protected separately using mutexes that belongs to the account itself. After accessing the specific account from the store, the account's mutex is acquired before read or mutation operations, and only unlocked when completed. This design allows clients to concurrently operate on different accounts while ensuring correctness on individual accounts.

For operations across two accounts (like `Transfer`), a naïvely locking each account could risk deadlock. To prevent this, we enforced a deterministic locking order.

```
1 // Lock both accounts in a consistent global order (lowest ID first)
2 // to prevent the "deadly embrace" deadlock scenario.
3 if fromAcc.AccountNumber < toAcc.AccountNumber {
4     fromAcc.Mu.Lock()
5     toAcc.Mu.Lock()
6 } else {
7     toAcc.Mu.Lock()
8     fromAcc.Mu.Lock()
9 }
10 defer fromAcc.Mu.Unlock()
11 defer toAcc.Mu.Unlock()
```

Listing 4: Deadlock-safe lock ordering in `Transfer` (`bank.go`)

This makes sure that any two goroutines performing a transfer between the same accounts acquire the locks in the same order. This eliminates circular waits.

4 Banking Services & Custom Operations

4.1 Core Services (S1–S4): Open, Close, Deposit, Withdraw

The server implements four core banking services:

OpenAccount (S1): Creates a new account with an initial balance, currency, and password. The server generates a unique account number, and stores the account under the global store write lock. This service returns the generated account number.

CloseAccount (S2): Deletes an account after verifying credentials. The handler retrieves the account, validates the user’s credentials, locks the account to prevent concurrent mutations, and removes it from the store. Notifications are sent to monitor subscribers.

Deposit (S3): Adds funds to an account. After credential validation, the handler acquires the account’s per-account mutex, checks that currency matches, increments the balance, and returns the new balance. The transaction is atomic with respect to other operations on the same account.

Withdraw (S4): Removes funds from an account. Similar to deposit, but additionally checks that sufficient balance exists. If any check fails (wrong credentials, currency mismatch, insufficient funds), an error status is returned without modifying state.

All modifying operations above trigger monitor callbacks after state changes. Errors are reported via the reply header’s status code, allowing clients to distinguish between errors like authentication failures and missing accounts.

All services use a shared set of status codes. They are transmitted as a `uint16` in the reply header’s status code field. Both the Go server’s (`protocol.go`) and the C++ client’s (`protocolStatus.h`) define the same enumerations to ensure consistency.

Table 2: Status Code Mapping (shared between Go server and C++ client)

Code	Name	Triggered by
0	SUCCESS	Operation completed without error
1	ACCOUNT_NOT_FOUND	Account number does not exist in the store
2	INVALID_CREDENTIALS	Password does not match the stored hash
3	ACCOUNT_MISMATCH	Holder name does not match the account
4	CURRENCY_MISMATCH	Request currency differs from the account’s currency
5	INSUFFICIENT_FUNDS	Withdrawal or transfer amount exceeds balance
6	TRANSFER_SAME_ACCOUNT	Source and destination accounts are identical
7	INTERNAL_SERVER_ERROR	Catch-all for unexpected server-side failures

On the server side, the handler maps errors from the banking service to these codes via `mapBankingError()` function. On the client side, the `to_string()` helper converts the received code into a readable label to output.

4.2 Monitor Service (S5)

The monitor feature allows clients to subscribe to real-time account update notifications. When a client issues a monitor request, specifying a timeout interval, the server registers the client’s address and expiration time within a mutex-protected subscriber map. A background goroutine periodically sweeps this map, removing expired entries.

Whenever a modifying operation (S1–S4) occurs, the handler invokes `NotifyAll()`, which marshals the account update once, then iterates through active subscribers (under lock) and sends a callback packet to each. Expired subscribers are lazily evicted during iteration. This design avoids holding account locks during network I/O.

On the client side, after receiving a monitor acknowledgment, the application enters a blocking receive

loop on the UDP socket, collecting callbacks until the timeout expires. User input is disabled during this period.

Figure 7 shows the monitor lifecycle. Client B will first register for monitoring. When Client A performs a deposit, the server will then push the callback to Client B in real time.




Figure 7: Sequence diagram of the monitor/callback lifecycle.

4.3 GetBalance (Idempotent Operation)

`GetBalance` queries an account's current balance without modifying state. After credential and currency validation, the server reads the balance under the account's mutex and replies with this value. Because the operation is read-only, it is inherently *idempotent*: repeating the request yields the same result without side effects. This property is valuable in UDP networks: if a reply is lost and the client retries, the outcome remains correct. Under both at-least-once and at-most-once semantics, idempotent operations are safe to retry, simplifying the reliability mechanisms.

4.4 Transfer (Non-Idempotent Operation)

`Transfer` moves funds atomically between two accounts. The request specifies source and destination account numbers, source credentials, amount, and currency. The handler's implementation must prevent two key issues: deadlock and inconsistency.

Deadlock avoidance: Transfer operations involve two accounts and face the risk of a deadlock, if they hold the resource that the other party requires, and wait indefinitely for it to be free. To avoid this, an ordering is imposed on account locks: the account with the smallest account number is locked first, then the other account is locked. This deterministic order ensures that no two simultaneous transfers will lock the same two accounts in different orders.

Atomic mutation: Under the locks, the handler verifies that both accounts exist, credentials are valid, currency matches, and the source has sufficient funds. If any check fails, both locks are released and an

error is returned. If valid, the handler subtracts from the source and adds to the destination atomically, then sends callbacks to monitor subscribers.

As Transfer modifies state (specifically, moves funds), it is *non-idempotent*: applying it twice would transfer funds twice, which is incorrect. This makes Transfer a critical case for duplicate suppression. Under at-least-once semantics, retransmitted requests would execute twice, doubling the transfer. Under at-most-once semantics, duplicate detection prevents re-execution, ensuring each transfer happens exactly once or not at all. This contrast motivates the need for correct invocation semantics, which is addressed in Section 5.

5 Fault Tolerance & Invocation Semantics

UDP is fundamentally an unreliable protocol, where packets could be lost, reordered, or duplicated. In this section, we cover two semantic models that tackle these issues, analyse their tradeoffs, and empirically validate them via a loss simulator.

5.1 At-Least-Once Semantics

In this model, during a failed transmission, the client retransmits requests with the same request ID, until a response arrives. The server executes each request it receives without filtering or caching of results.

Client-side mechanism: After encoding and serialising a request, the client sends it via UDP, and awaits a reply with a 5-second timeout. If no reply arrives within the timeout, the client retransmits with an exponential backoff, doubling the timeout period. The timeout sequence is 5 s → 10 s → 20 s → 40 s → 80 s. This gives the network up to 155 seconds in total to deliver a reply. This occurs up to a max of 5 attempts.

Server-side mechanism: The dispatcher receives the request header and immediately dispatches it to the handler for execution, without filtering for duplication. All requests that are received are processed regardless of whether a similar request had previously arrived.

In the event that the request is lost, and the server never received it, when the client recovers and retransmits within 3 attempts, the server will process the request at least once. However, if the retransmission by the client is caused by a lost reply, the server may execute the same operation multiple times.




Figure 8: At-least-once semantics: retransmission on timeout can cause duplicate execution.

As depicted in Figure 8, at-least-once is safe only for idempotent operations like GetBalance, which can be repeated and still yield the same result. However, non-idempotent operations like Transfer are not suitable for this, as duplicate executions can lead to incorrect behaviour. For instance, a duplicate transfer would unexpectedly move funds twice, making this semantic unsuitable for banking operations without additional protections.

5.2 At-Most-Once Semantics

For at most once, the client has a similar retry mechanism as at-least-once mode. However, the server detects and suppresses duplicates by caching replies, ensuring that each operation is executed at most one time, or zero times if all attempts are lost before execution.

Client-side mechanism: While the retry mechanism is similar to at-least-once, the client increments the request ID within the message for each session. This allows the server to distinguish fresh requests from retransmissions.

Server-side mechanism: The dispatcher maintains a request history cache, which is a two-dimensional

map that is keyed firstly by the client's IP address and port, and secondly by the request ID. Before dispatching a request to the handler, it will parse the message header and extract both keys. It will check the cache for a match:

1. If a match is found, the request is a retransmission, and we return the cached reply immediately.
2. Else, the handler executes and caches the reply bytes before returning them.

This is the core duplication detection logic in the dispatcher:

```

1 // Check cache before executing --- a hit means the client
2 // is retransmitting because our previous reply got lost.
3 if cached, found := d.history.Lookup(clientKey, header.RequestID); found {
4     log.Printf("duplicate detected for request %d, returning cached reply",
5         header.RequestID)
6     return cached, nil
7 }
8 // First time seeing this request: execute and cache
9 reply := d.handler(data, addr)
10 d.history.Store(clientKey, header.RequestID, reply)

```

Listing 5: Duplicate detection in the at-most-once dispatcher (dispatcher.go)




Figure 9: At-most-once semantics.

As shown in Figure 9, each request is processed a maximum of once. This is the ideal for banking, where monetary transfers must execute exactly once per submission as intended. However, this semantic requires the system to scale its memory usage proportionally to the number of active clients and their request backlog, which could grow greatly in production. Additionally, the use of an incremented request ID in the client introduces coordination overhead in this implementation.

To address these issues, we structured the reply history to use a two level map. The outer key is client's IP:port string and the inner key is the 32-bit request ID. This structure will allow for $O(1)$ lookup time per request. When the server receives a new request with an ID N , the dispatcher calls `EvictBefore(clientAddr, N - windowSize)`. This would remove all cache entries for request IDs less than $N - 10$, bounding cache memory to $O(\text{clients} \times \text{windowSize})$.

5.3 Loss Simulator Design

To validate the difference between both semantics when messages are lost, a loss simulator (`losssim.go`) was implemented in the server. It wraps a `rand.Rand` instance with a fixed seed, allowing for experiments to be reproducible across runs (two runs with the same seed produce the same drop sequences). The drop probability is configurable via the CLI argument (i.e. `./server at-most-once 0.3` for 30% loss). It is configured with a probability between 0.0 and 1.0.

The loss simulator can be injected at either of two points:

1. **Request-level:** After the server's UDP socket receives a datagram, the loss simulator may discard it before it is parsed. This forces the client to timeout and retransmit.
2. **Reply-level loss:** After the handler produces a reply, the loss simulator may drop it before it is sent to the client, simulating the scenario where the client's receive mechanism times out.

For this implementation, we focus on reply-level loss only. This is the more dangerous scenario whereby the server has already executed and mutated the state. The client, having not received the confirmation, would retransmit and trigger a duplicate execution under the at-least-once semantics. This is the failure that we are trying to demonstrate.

5.4 Experimental Validation

To validate the semantic models, a suite of 38 integration tests was developed. These tests launch the real C++ client and Go server as separate processes, communicating over localhost UDP. Key test scenarios include:

Table 3: Experimental validation: expected vs. observed outcomes.

Operation	Semantics	Loss	Expected Outcome	Observed Result
GetBalance	at-least-once	0%	Correct balance returned	Balance = \$1000 ✓
GetBalance	at-least-once	30%	Correct (idempotent, safe to retry)	Balance = \$1000 ✓
Deposit \$100	at-least-once	0%	Balance +\$100	Balance = \$1100 ✓
Deposit \$100	at-least-once	30%	Duplicate: +\$100 to +\$200	Balance = \$1200 ×
Deposit \$100	at-most-once	30%	Exactly +\$100 (dedup)	Balance = \$1100 ✓
Transfer \$50	at-least-once	30%	Duplicate: source −\$100	Alice=\$900, Bob=\$600 ×
Transfer \$50	at-most-once	30%	Exactly −\$50 (dedup)	Alice=\$950, Bob=\$550 ✓

5.5 Results & Discussion

The integration test suite confirms the theoretical predictions.

Firstly, for idempotent operations like GetBalance, the correct balance is returned under both semantics and all loss rates. The operation is inherently safe to retry because it does not modify any states. Thus, at-least-once is acceptable for GetBalance.

Secondly, for non-idempotent operations (Deposit, Withdraw, Transfer), we analyse at-least-once semantics. When simulated packet loss occurs, duplicate requests cause multiple executions, which is unsafe for non-idempotent banking operations. For example:

- Transfer \$50 with 30% loss rate: If a single response packet is lost, the source account is debited twice, causing a \$100 total decrease instead of \$50.

For non-idempotent operations under at-most-once semantics, the server's request history cache detects retransmissions and prevents duplicate executions, which is safe for banking operations. In this scenario, we can validate that the response is returned from the cache:

- Transfer \$50 with 30% loss rate: Even if multiple response packets were lost, source is debited exactly \$50 and destination is credited exactly \$50.

Our results are tested against an integration test suite. It runs the real C++ Client against the real Go Server in our CI Pipeline. The following is an example test output from our TestComparison_SameScenarioBothModes. It executes the same Transfer scenario of both semantics but at 30% packet loss.

Listing 6: Integration test output: at-least-once Transfer under 30% loss

```

=== RUN    TestComparison_SameScenarioBothModes/at-least-once
[Dispatcher] At-least-once: executed request 5 from 127.0.0.1:54321
[Dispatcher] At-least-once: executed request 5 from 127.0.0.1:54321
integration_test.go:1285: [at-least-once] Alice=800.00, Bob=700.00
integration_test.go:1253: REPORT EVIDENCE: Transfer duplicated!
--- PASS

```

Listing 7: Integration test output: at-most-once Transfer under 30% loss

```

=== RUN    TestComparison_SameScenarioBothModes/at-most-once
[Dispatcher] At-most-once: executed and cached request 5 from 127.0.0.1:54321
[Dispatcher] At-most-once: duplicate detected for request 5 from 127.0.0.1:54321,
                returning cached reply
    integration_test.go:1285: [at-most-once] Alice=900.00, Bob=600.00
--- PASS

```

In the at-least-once run (Listing 6), request 5 (the Transfer) was executed twice by the server because the first reply was dropped and the client retransmitted. Alice lost \$200 instead of \$100. In the at-most-once run (Listing 7), the server’s dispatcher detected the duplicate retransmission and returned the cached reply without re-executing the handler, preserving the correct balances.

To conclude, the choice of invocation semantics directly impacts correctness of the system. For banking systems that require this accuracy, and handles non-idempotent operations, at-most-once semantics are necessary.

6 Testing Strategy

The project employs a two-tier testing approach: unit tests for individual components with mocked dependencies, and end-to-end integration tests exercising the full system.

6.1 Unit Tests (C++ Client)

The client includes 48 unit tests across four GoogleTest suites, with dependency mocks to isolate each layer.

Table 4: C++ client unit test suites and their scope.

Suite	Scope	Key checks
CommandEncoder	TLV request/response encoding for all services	Big-endian fields, variable-length strings, and IEEE-754 amount round-trip correctness
MessageSerializer	18-byte message header and payload framing	Header round-trip (IPv4, port, status, length), short-packet rejection, overflow-safe length handling
UDPSocket	UDP send/receive integration on loopback	Timeout options (SO_RCVTIMEO, SO_SNDTIMEO), SO_REUSEADDR, and getsockname() correctness
BankClient	CLI parsing and command construction	Input validation (account/currency/password) and correct request assembly before dispatch

By mocking components like BaseSocket and BaseCommandEncoder, the tests do not rely on network dependencies, and execute deterministically.

6.2 Integration Tests (End-to-End)

The integration test suite (integration_test.go) is written in Go and runs 38 test cases that pass on both Windows and Mac. Each test:

1. Spawns the real Go server with specified flags (e.g., ‘-semantics=at-most-once -loss-rate=0.3’).
2. Spawns the real C++ client, instructing it to connect to the server’s UDP address.
3. Executes a sequence of operations by injecting CLI commands via pipes.
4. Validates outcomes by querying the server’s in-memory state or comparing received replies.
5. Terminates both processes and reports pass/fail.

Table 5: Integration test categories and coverage summary.

Category	Coverage summary
Baseline (no loss)	Verifies all seven services end-to-end under normal network conditions.
Concurrency	Runs multiple clients in parallel to validate consistency and account-total conservation.
Fault tolerance (10%/30%/50%)	Compares at-least-once vs at-most-once behavior under packet loss, and confirms that duplicate side effects only occur in at-least-once for non-idempotent operations.
Stress and edge cases	Tests rapid-fire requests, large balances, and invalid boundary inputs (e.g., withdraw \$0).
Isolation	Executes each test in an isolated process/port with fresh server state to avoid cross-test interference.

7 Assumptions and Limitations

7.1 System Assumptions

Our implementation has two main assumptions. First, each account is fixed to a single currency (SGD, USD, or EUR). This means that transfers between accounts of different currencies are rejected. Second, all account data is stored in a transient Go map with no persistence. If the server crashes, all data is lost.

7.2 Known Limitations

Our system also has several limitations. The at-most-once request history cache grows with each unique request. This could exhaust server memory over time. While our current implementation supports manual eviction via a sliding window, a time based/size based eviction policy would be better. On the security front, all messages are transmitted in plaintext over UDP. Account numbers, balances, and passwords by right should be encrypted. Our protocol enforces a maximum 8-byte password field, which truncates longer passwords.

8 Conclusion

This project set out to consolidate our understanding of interprocess communication and remote invocation. It successfully gave us real world understanding for which reading the textbook alone could not. When we were building the wire protocol by hand, deciding on header layouts, TLV field IDs, endianness conventions, we gained a more concrete appreciation for the tedious yet crucial details that RPC frameworks abstract away for us. Every misaligned byte we had to debug, every silent corruption we had to fix taught us the importance of marshalling standards.

The most important lesson came about when working on both invocation semantics side by side. It was one thing to know the theory but implementing it was a whole different ball game. We unfortunately watched the Transfer debit an account twice in test runs way too many times. This experience gave us a deeper appreciation for the request history cache, making its memory cost feel less like a textbook trade off and more like a necessary engineering decision.

Our choice of working with two languages was also really eye opening. It forced us to confront many of the assumptions we would never have noticed in a single language system. Go's fixed [8]byte password type and C++'s variable-length `std::string` only agreed on the wire because we explicitly defined the contract. This is the same coordination problem that any distributed team faces when integrating any system.

In the end, this was an extremely fruitful project. It reinforced the principle that in a distributed system, correctness is not handled by application logic alone. We often take for granted the protocols, semantics, and failure models upon which we heavily depend.